

Proximal Algorithms

For distributed and constrained optimization
Niru Maheswaranathan
October 14, 2015

Outline

- Distributed optimization in machine learning
- Dual ascent and augmented Lagrangians
- Proximal Operators and Proximal Algorithms
- Evaluating proximal operators
- Interactive Demo

Outline

- Distributed optimization in machine learning
- Dual ascent and augmented Lagrangians
- Proximal Operators and Proximal Algorithms
- Evaluating proximal operators
- Interactive Demo


Distributed optimization in machine learning

We would like to solve the following:

$$\text{minimize} \quad \sum_i f_i(x)$$


where i indexes data batches, different loss functions, regularizers, etc.

Ideally, we would like to distribute the work amongst many workers


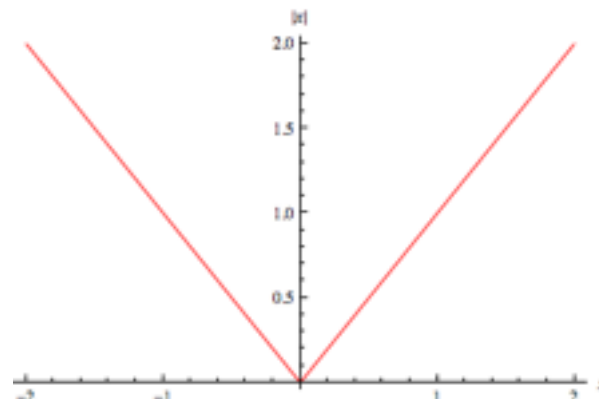


And be able to handle: constraints, non-smooth terms, poor conditioning

$$f_i(x) = \mathbf{I}\{x \in \mathcal{C}\}$$



$$f_i(x) = \lambda \|x\|_1$$



Gradient Descent

$$\text{minimize } \sum_i f_i(x)$$


Gradient descent
update:

$$x^{k+1} = x^k - \alpha^k \sum_i \nabla f_i(x)$$


- ☒ Works for smooth, convex functions
- ☒ Simple implementation
- ☒ Can distribute the gradient computation for each subproblem
- ☐ Constraints? Non-smooth functions? Ill-conditioned problems?

Gradient descent on a non-smooth function

Non-smooth function



Smooth function



Outline

- Distributed optimization in machine learning
- Dual ascent and augmented Lagrangians
- Proximal Operators and Proximal Algorithms
- Evaluating proximal operators
- Interactive Demo

Dual Ascent

Consider:

$$\begin{array}{ll} \text{minimize} & f(x) \\ \text{subject to} & Ax = b \end{array}$$

Lagrangian:

$$L(x, \lambda) = f(x) + \lambda^T (Ax - b)$$

Decomposes over x

Dual:

$$g(\lambda) = \min_x f(x) + \lambda^T (Ax - b)$$

Dual ascent:

$$x^{k+1} = \operatorname{argmin}_x L(x, \lambda^k)$$

$$\lambda^{k+1} = \lambda^k + \alpha(Ax^{k+1} - b)$$

- * **Brittle; only converges under strong assumptions**
- * **Decomposes across subproblems!**

Augmented Lagrange Methods

(also known as: Method of Multipliers)

$$\tilde{L}(x, \lambda) = f(x) + \lambda^T (Ax - b) + \frac{\rho}{2} \|Ax - b\|^2$$

- 😊 Adds minimum curvature to $f(x)$
- 😞 Not decomposable over x



$$x^{k+1} = \underset{x}{\operatorname{argmin}} \tilde{L}(x, \lambda^k)$$
$$\lambda^{k+1} = \lambda^k + \rho(Ax^{k+1} - b)$$

* **Converges under weak assumptions!**

* **But we lose decomposability :(**

Alternating Direction Method of Multipliers (ADMM)

$$\begin{array}{ll}\text{minimize} & f(x) + g(z) \\ \text{subject to} & x - z = 0\end{array}$$

Augmented Lagrangian:

$$\tilde{L}(x, z, \lambda) = f(x) + g(z) + \lambda^T(x - z) + \frac{\rho}{2} \|x - z\|_2^2$$


ADMM:

$$\begin{aligned}x^{k+1} &= \underset{x}{\operatorname{argmin}} \tilde{L}(x, z^k, \lambda^k) \\ z^{k+1} &= \underset{z}{\operatorname{argmin}} \tilde{L}(x^{k+1}, z, \lambda^k) \\ \lambda^{k+1} &= \lambda^k + \rho(x^{k+1} - z^{k+1})\end{aligned}$$

- * **Converges under weak assumptions!**
- * **And it is easily parallelizable! (as we will see...)**

Scaled form of ADMM

$$x^{k+1} = \operatorname{argmin}_x \tilde{L}(x, z^k, \lambda^k)$$

$$\tilde{L}(x, z, \lambda) = f(x) + g(z) + \lambda^T (x - z) + \frac{\rho}{2} \|x - z\|_2^2$$


(pull the linear term inside the quadratic)

Scaled form:

$$\tilde{L}(x, z, \lambda) = f(x) + g(z) + \frac{\rho}{2} \|x - z + u\|_2^2 + \text{const.}$$

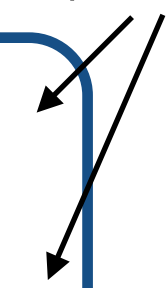
(where $u = \lambda/\rho$)

These are *proximal operators*

ADMM
(scaled form)

$$x^{k+1} = \operatorname{argmin}_x f(x) + \frac{\rho}{2} \|x - z^k + u^k\|_2^2$$

$$z^{k+1} = \operatorname{argmin}_z g(z) + \frac{\rho}{2} \|x^{k+1} - z + u^k\|_2^2$$

$$u^{k+1} = u^k + x^{k+1} - z^{k+1}$$


Outline

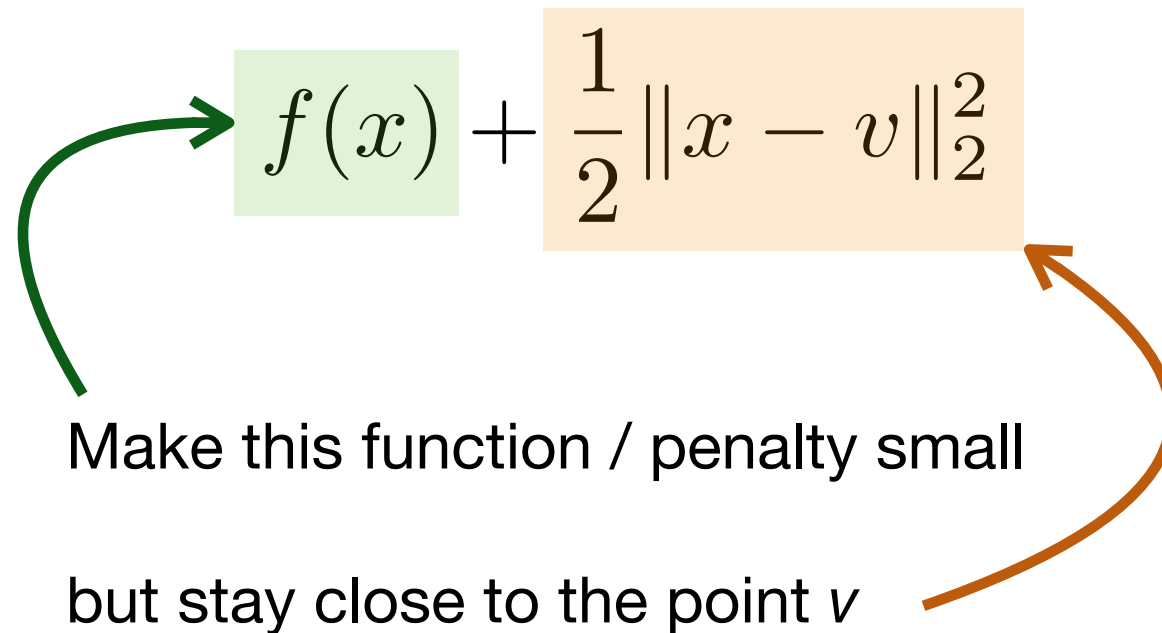
- Distributed optimization in machine learning
- Dual ascent and augmented Lagrangians
- Proximal Operators and Proximal Algorithms
- Evaluating proximal operators
- Interactive Demo

Proximal operators

Definition


$$\text{prox}_f(v) = \operatorname{argmin}_x \left(f(x) + \frac{1}{2} \|x - v\|_2^2 \right)$$

Minimize



Proximal operator for the l1-norm

$$f(x) + \frac{\rho}{2} \|x - v\|^2$$



Interpretations

$$\text{prox}_{\rho, f}(v) = \underset{x}{\operatorname{argmin}} \left(f(x) + \frac{\rho}{2} \|x - v\|_2^2 \right)$$

1. Function smoothing

- The added term *smooths* out the function $f(x)$ locally near v

2. MAP estimation with a Gaussian prior

- Think of $f(x)$ as a negative log-likelihood term
- The added term is the negative log-probability for a Gaussian prior centered at v with identity covariance

3. Connection to gradient methods

Aside: proximal operators and gradient descent

First-order approximation to the function near the point v :

$$\hat{f}(x) \approx f(v) + \nabla f(v)^T (v - x)$$

Let's apply the proximal operator to this approximation:

$$\text{prox}_f(x) = \underset{x}{\operatorname{argmin}} \left(\hat{f}(x) + \frac{1}{\alpha} \|x - v\|^2 \right)$$

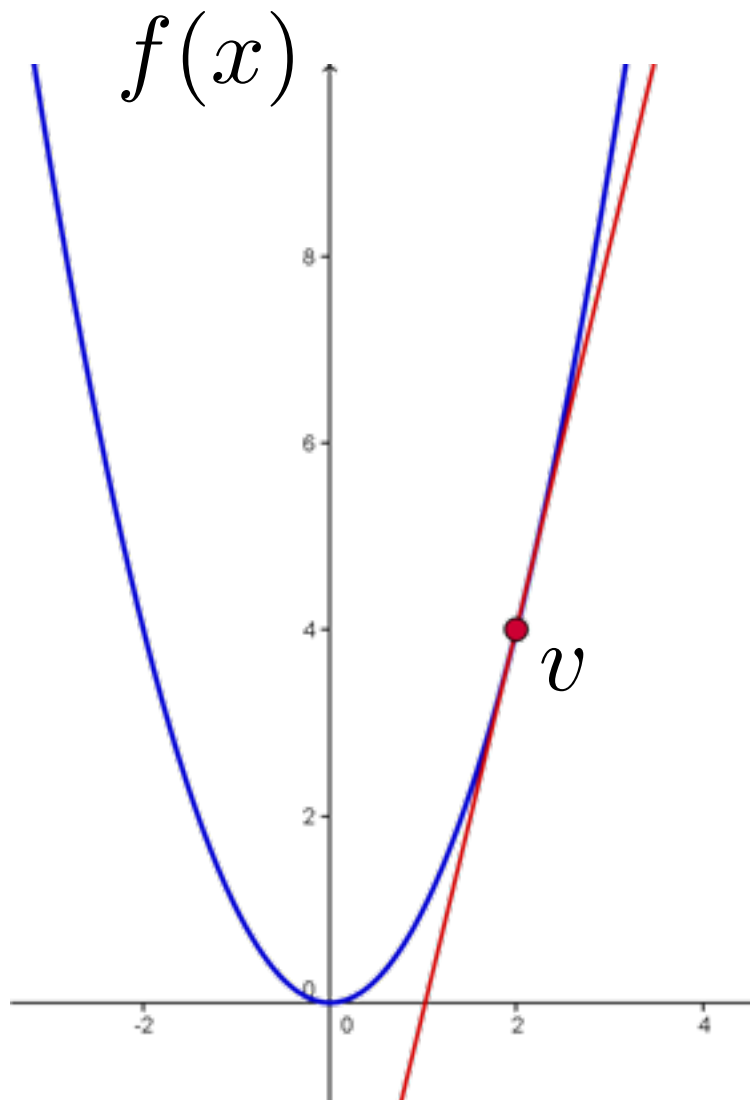
$$\underset{x}{\operatorname{argmin}} f(v) + \nabla f(v)^T (v - x) + \frac{1}{\alpha} \|x - v\|^2$$

Analytic minimization:

$$\nabla \hat{f}(x) = \nabla f(v) + \frac{1}{\alpha} (x - v) = 0$$

$$x = v - \alpha \nabla f(v)$$

yields the gradient descent update equation



Proximal algorithms (recap)

Problem: minimize $f(x) + g(x)$

Operator splitting:

 minimize $f(x) + g(z)$
 subject to $x = z$

Algorithm (ADMM):

$$\begin{aligned}x^{k+1} &:= \mathbf{prox}_{\lambda f}(z^k - u^k) \\z^{k+1} &:= \mathbf{prox}_{\lambda g}(x^{k+1} + u^k) \\u^{k+1} &:= u^k + x^{k+1} - z^{k+1},\end{aligned}$$

*We benefit if proximal operators for $f(x)$ and $g(x)$ are easy to evaluate,
but the operator for $f+g$ is difficult to evaluate*

Proximal consensus

Problem: minimize $\sum_i f_i(x)$

Algorithm:

$$x_i^{k+1} = \operatorname{argmin}_x \left(f_i(x) + \frac{\rho}{2} \|x_i - \bar{x}^k + u_i^k\|_2^2 \right)$$

$$\bar{x}^{k+1} = \frac{1}{m} \sum_{i=1}^m x_i^{k+1}$$

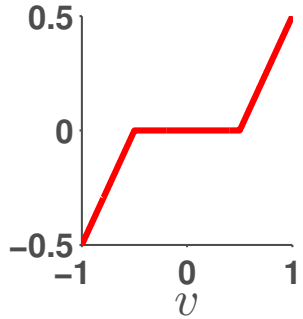
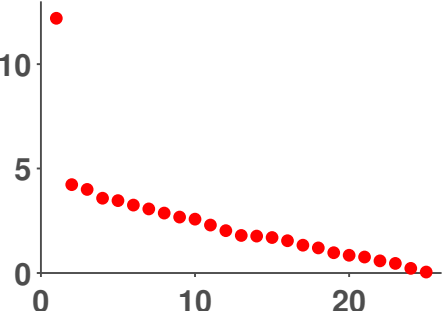
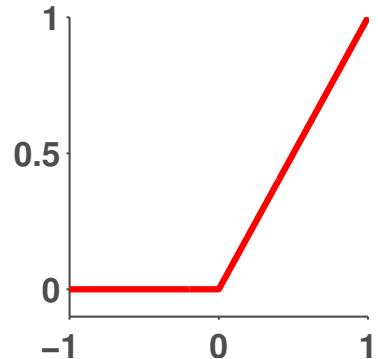
$$u_i^{k+1} = u_i^k + (x_i^{k+1} - \bar{x}^{k+1})$$

- ☑ Works for smooth, non-smooth, ill-conditioned convex functions
- ☑ Can run the proximal update for each subproblem in parallel
- ☑ Guaranteed to converge, finds reasonable solutions quickly (similar to first order methods)

Outline

- Distributed optimization in machine learning
- Dual ascent and augmented Lagrangians
- Proximal Operators and Proximal Algorithms
- Evaluating proximal operators
- Interactive Demo

Evaluating proximal operators

Property	$f(x)$	$\text{prox}_f(v)$
Least squares	$\frac{1}{2} \ Ax - b\ _2^2$	$(A + \rho I)^{-1}(\rho v - b)$
Sparsity (l1-norm)	$\ x\ _1$	<p>Soft thresholding</p> 
Low-rank matrix (nuclear norm)	$\ X\ _*$	<p>Soft thresholding the singular values</p> 
Non-negative values	$\mathbf{I}\{x \geq 0\}$	<p>Clip negative values</p> 

Outline

- Distributed optimization in machine learning
- Dual ascent and augmented Lagrangians
- Proximal Operators and Proximal Algorithms
- Evaluating proximal operators
- Interactive Demo

Demo

<https://github.com/nirum/descent-tutorial>

References

- N. Parikh and S. Boyd, *Proximal Algorithms*, 2014. http://stanford.edu/~boyd/papers/pdf/prox_algs.pdf
- S. Boyd, N. Parikh, E. Chu, B. Peleato, and J. Eckstein, *Distributed Optimization and Statistical Learning via the Alternating Direction Method of Multipliers*, 2011. http://stanford.edu/~boyd/papers/pdf/admm_distr_stats.pdf